

exploration

January 21, 2026

```
[ ]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, PolynomialFeatures
from sklearn.linear_model import LinearRegression, Ridge, Lasso
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score
```

```
[ ]: from sklearn.datasets import fetch_openml
boston = fetch_openml(name='boston', version=1, parser='auto')
X = pd.DataFrame(boston.data, columns=boston.feature_names)
y = pd.Series(boston.target, name='MEDV')

print(f"Dataset shape: {X.shape}")
print(f"Target shape: {y.shape}")
print(f"\nFeatures: {list(X.columns)}")
```

Dataset shape: (506, 13)

Target shape: (506,)

Features: ['CRIM', 'ZN', 'INDUS', 'CHAS', 'NOX', 'RM', 'AGE', 'DIS', 'RAD',
'TAX', 'PTRATIO', 'B', 'LSTAT']

```
[ ]: missing_values = X.isnull().sum()
print(f"Missing values per feature:\n{missing_values}")
print(f"Total missing values: {missing_values.sum()}")
```

Missing values per feature:

CRIM	0
ZN	0
INDUS	0
CHAS	0
NOX	0
RM	0
AGE	0
DIS	0
RAD	0
TAX	0

```
PTRATIO      0  
B            0  
LSTAT       0  
dtype: int64  
Total missing values: 0
```

```
[ ]: #EDA  
print("\nDataset Statistics:")  
print(X.describe())  
print(f"\nTarget Variable Statistics:")  
print(y.describe())
```

Dataset Statistics:

	CRIM	ZN	INDUS	NOX	RM	AGE	\
count	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	
mean	3.613524	11.363636	11.136779	0.554695	6.284634	68.574901	
std	8.601545	23.322453	6.860353	0.115878	0.702617	28.148861	
min	0.006320	0.000000	0.460000	0.385000	3.561000	2.900000	
25%	0.082045	0.000000	5.190000	0.449000	5.885500	45.025000	
50%	0.256510	0.000000	9.690000	0.538000	6.208500	77.500000	
75%	3.677083	12.500000	18.100000	0.624000	6.623500	94.075000	
max	88.976200	100.000000	27.740000	0.871000	8.780000	100.000000	

	DIS	TAX	PTRATIO	B	LSTAT
count	506.000000	506.000000	506.000000	506.000000	506.000000
mean	3.795043	408.237154	18.455534	356.674032	12.653063
std	2.105710	168.537116	2.164946	91.294864	7.141062
min	1.129600	187.000000	12.600000	0.320000	1.730000
25%	2.100175	279.000000	17.400000	375.377500	6.950000
50%	3.207450	330.000000	19.050000	391.440000	11.360000
75%	5.188425	666.000000	20.200000	396.225000	16.955000
max	12.126500	711.000000	22.000000	396.900000	37.970000

Target Variable Statistics:

```
count      506.000000  
mean      22.532806  
std       9.197104  
min      5.000000  
25%     17.025000  
50%     21.200000  
75%     25.000000  
max      50.000000  
Name: MEDV, dtype: float64
```

```
[ ]: scaler = StandardScaler()  
X_train_scaled = scaler.fit_transform(X_train)  
X_test_scaled = scaler.transform(X_test)
```

```

# Convert back to DataFrame for readability
X_train_scaled = pd.DataFrame(X_train_scaled, columns=X.columns)
X_test_scaled = pd.DataFrame(X_test_scaled, columns=X.columns)

print("Features normalized (mean=0, std=1)")
print(f"\nSample of normalized features:")
print(X_train_scaled.head())

```

Features normalized (mean=0, std=1)

Sample of normalized features:

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	\
0	1.287702	-0.500320	1.033237	-0.278089	0.489252	-1.428069	1.028015	
1	-0.336384	-0.500320	-0.413160	-0.278089	-0.157233	-0.680087	-0.431199	
2	-0.403253	1.013271	-0.715218	-0.278089	-1.008723	-0.402063	-1.618599	
3	0.388230	-0.500320	1.033237	-0.278089	0.489252	-0.300450	0.591681	
4	-0.325282	-0.500320	-0.413160	-0.278089	-0.157233	-0.831094	0.033747	
	DIS	RAD	TAX	PTRATIO	B	LSTAT		
0	-0.802173	1.706891	1.578434	0.845343	-0.074337	1.753505		
1	0.324349	-0.624360	-0.584648	1.204741	0.430184	-0.561474		
2	1.330697	-0.974048	-0.602724	-0.637176	0.065297	-0.651595		
3	-0.839240	1.706891	1.578434	0.845343	-3.868193	1.525387		
4	-0.005494	-0.624360	-0.584648	1.204741	0.379119	-0.165787		

```

[ ]: fig, axes = plt.subplots(2, 2, figsize=(15, 10))

# Distribution of target variable
axes[0, 0].hist(y, bins=30, edgecolor='black', alpha=0.7)
axes[0, 0].set_title('Distribution of House Prices (MEDV)', fontsize=12, u
    ↪fontweight='bold')
axes[0, 0].set_xlabel('Price ($1000s)')
axes[0, 0].set_ylabel('Frequency')

# Correlation heatmap (top features)
correlation_matrix = pd.concat([X, y], axis=1).corr()
top_features = correlation_matrix['MEDV'].abs().sort_values(ascending=False)[1:
    ↪8].index
sns.heatmap(pd.concat([X[top_features], y], axis=1).corr(), annot=True, u
    ↪cmap='coolwarm',
            center=0, ax=axes[0, 1], fmt=' .2f')
axes[0, 1].set_title('Correlation Heatmap (Top Features)', fontsize=12, u
    ↪fontweight='bold')

# Scatter plot: RM vs MEDV
axes[1, 0].scatter(X['RM'], y, alpha=0.5, edgecolors='k', linewidth=0.5)

```

```


axes[1, 0].set_title('Rooms vs House Price', fontsize=12, fontweight='bold')
axes[1, 0].set_xlabel('Average Number of Rooms (RM)')
axes[1, 0].set_ylabel('House Price ($1000s)')

# Scatter plot: LSTAT vs MEDV
axes[1, 1].scatter(X['LSTAT'], y, alpha=0.5, edgecolors='k', linewidth=0.5, color='coral')
axes[1, 1].set_title('LSTAT vs House Price', fontsize=12, fontweight='bold')
axes[1, 1].set_xlabel('% Lower Status Population (LSTAT)')
axes[1, 1].set_ylabel('Median House Price ($1000s)')

plt.tight_layout()
plt.show()

# 4.3 Feature Correlations with Target
print("\nFeature Correlations with House Price:")
correlations = pd.concat([X, y], axis=1).corr()['MEDV'].
    ↪sort_values(ascending=False)
print(correlations)

```



Feature Correlations with House Price:
 MEDV 1.000000

```
RM           0.695360
ZN           0.360445
B            0.333461
DIS          0.249929
CHAS         0.175260
AGE          -0.376955
RAD           -0.381626
CRIM         -0.388305
NOX          -0.427321
TAX           -0.468536
INDUS        -0.483725
PTRATIO      -0.507787
LSTAT        -0.737663
Name: MEDV, dtype: float64
```

```
[ ]: #Train/test split
print("\nSplitting data into train/test sets (80/20)")
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
                                                    random_state=42)

print(f"Training set size: {X_train.shape[0]} samples")
print(f"Test set size: {X_test.shape[0]} samples")
print(f"Train/Test ratio: {X_train.shape[0]}/{X_test.shape[0]}:.1f]:1")
```

```
Splitting data into train/test sets (80/20)
Training set size: 404 samples
Test set size: 102 samples
Train/Test ratio: 4.0:1
```